



# 人工智慧作業報告簡報

信息工程學院 2021 級 干皓丞 2101212850

Planning in Markov Decision Processes with Gap-Dependent Sample Complexity

用間隙相關樣本複雜性來做出利用馬爾可夫決策過程所做出的規劃

<https://arxiv.org/abs/2006.05879>

<https://proceedings.neurips.cc/paper/2020>

<https://proceedings.neurips.cc/paper/2020/hash/0d85eb24e2add96ff1a7021f83c1abc9-Abstract.html>

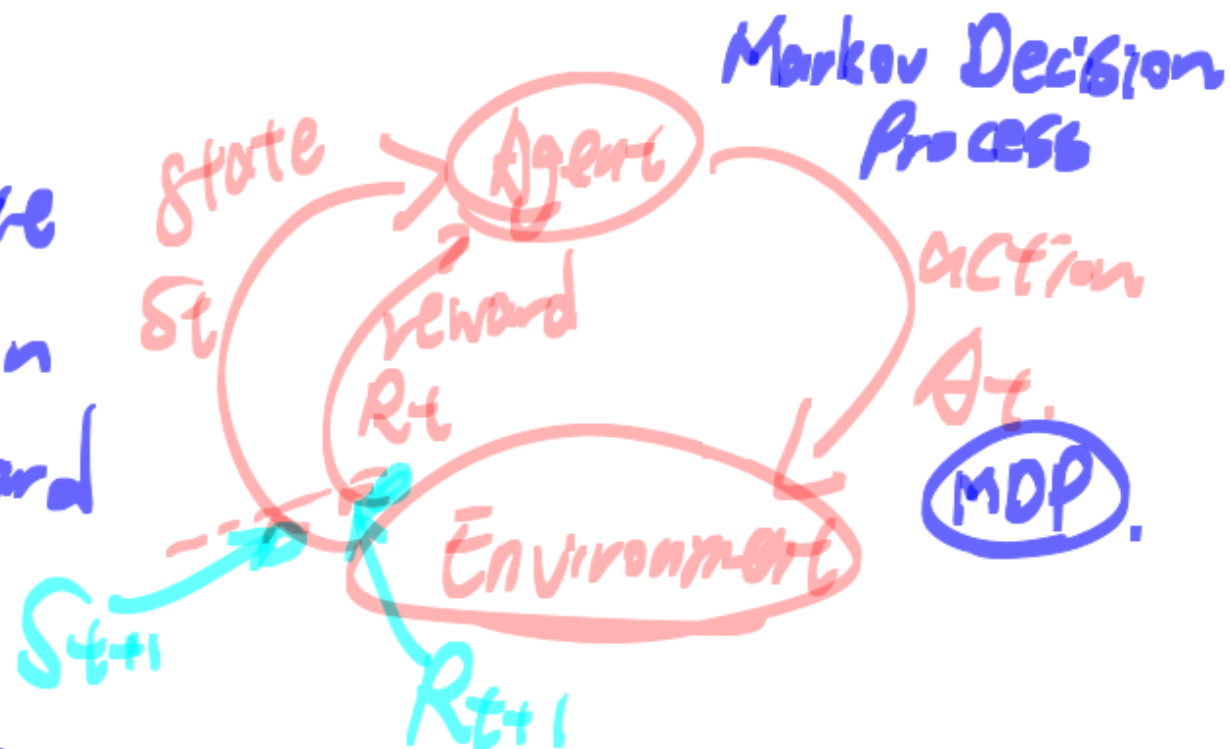
Part of Advances in Neural Information Processing Systems 33 (NeurIPS 2020)





# Motivation 动机

3. 代理學習  
3. 要素  
2. 特性  
1. 所有過程都是隨機的。  
4. 下一階段狀態轉移是隨機的。  
現在的水題百問



狀態 State,  $S = \{s_0, s_1, \dots, s_n\}$   
動作 Action,  $A(s) = \{a_0, a_1, \dots, a_m\}$   
獎勵 Reward,  $R = \{r_0, r_1, \dots, r_n\}$   
轉移概率 Dynamic,  $P = p(s', r | s, a)$   
for all  $s \in S, a \in A$   
 $M = (S, A, P, r)$

該研究的研究者提出名為一個名為 MDP-GapE 的算法，這個演算法是一種新的基於“軌跡的蒙特卡洛樹搜索演算法”(trajectory-based Monte-Carlo Tree Search algorithm)，此算法用於在馬爾可夫決策過程中進行規劃，該算法會在過程中進行規畫，且而這個過程的轉換是具有‘有限支持’(a finite support)。



# Intuition 直觉

Table 1: Different settings of planning algorithms in the literature

Setting	Input	Output	Optimality criterion
(1) Fixed confidence (action-based)	$\varepsilon, \delta$	$\hat{a}_n$	$\mathbb{P}(\bar{r}_n(\hat{a}_n) \leq \varepsilon) \geq 1 - \delta$
(2) Fixed confidence (value-based)	$\varepsilon, \delta$	$\hat{V}(s_1)$	$\mathbb{P}\left( \hat{V}(s_1) - V^*(s_1)  \leq \varepsilon\right) \geq 1 - \delta$
(3) Fixed budget	$n$ (budget)	$\hat{a}_n$	$\mathbb{E}[\bar{r}_n(\hat{a}_n)]$ decreasing in $n$
(4) Anytime	-	$\hat{a}_n$	$\mathbb{E}[\bar{r}_n(\hat{a}_n)]$ decreasing in $n$

We propose an algorithm in the *fixed confidence setting*  $(\varepsilon, \delta)$ : after  $n$  calls to the generative model, the algorithm should return an action  $\hat{a}_n$  such that  $\bar{r}_n(\hat{a}_n) \leq \varepsilon$  with probability at least  $1 - \delta$ . We prove that its *sample complexity*  $n$  is bounded in high probability by a quantity that depends on

the sub-optimality gaps of the actions that are applicable in state  $s_1$ . We also provide experiments showing its effectiveness. The only assumption that we make on the MDP is that the support of the transition probabilities  $p_h(\cdot|s, a)$  should have cardinality bounded by  $B < \infty$ , for all  $s, a$  and  $h$ .



# Justification 理由

Table 2: Algorithms with sample complexity guarantees

Algorithm	Setting	Sample complexity	Remarks
Sparse Sampling [19]	(1)-(2)	$H^5(BK)^H / \varepsilon^2$ or $\varepsilon^{-\left(2 + \frac{\log(K)}{\log(1/\gamma)}\right)}$	proved in Lemma 1
OLOP [2]	(3)	$\varepsilon^{-\max\left(2, \frac{\log \kappa}{\log(1/\gamma)}\right)}$	open loop, $\kappa \in [1, K]$
OP [3]	(4)	$\varepsilon^{-\frac{\log \kappa}{\log(1/\gamma)}}$	known MDP, $\kappa \in [0, BK]$
BRUE [8]	(4)	$H^4(BK)^H / \Delta^2$	minimal gap $\Delta$
StOP [27]	(1)	$\varepsilon^{-\left(2 + \frac{\log \kappa}{\log(1/\gamma)} + o(1)\right)}$	$\kappa \in [0, BK]$
TrailBlazer [13]	(2)	$\varepsilon^{-\max\left(2, \frac{\log(B\kappa)}{\log(1/\gamma)} + o(1)\right)}$	$\kappa \in [1, K]$
SmoothCruiser [14]	(2)	$\varepsilon^{-4}$	only regularized MDPs
MDP-GapE (ours)	(1)	$\sum_{a_1 \in \mathcal{A}} \frac{H^2(BK)^{H-1} B}{(\Delta_1(s_1, a_1) \vee \Delta \vee \varepsilon)^2}$	see Corollary 1

**Corollary 1.** *The number of episodes used by MDP-GapE satisfies*

$$\mathbb{P} \left( \tau = \mathcal{O} \left( \sum_{a_1} \frac{(BK)^{H-1}}{(\Delta_1(s_1, a_1) \vee \Delta \vee \varepsilon)^2} \left[ \log \left( \frac{1}{\delta} \right) + BH \log(BK) \right] \right) \right) \geq 1 - \delta .$$



# Framework 框架

---

**Algorithm 1** MDP-GapE

---

```
1: Input: confidence level  $\delta$ , tolerance  $\varepsilon$ 
2: initialize data lists  $\mathcal{D}_h \leftarrow []$  for all  $h \in [H]$ 
3: for  $t = 1 \dots$  do
4:   //Update confidence bounds
5:    $U_h^{t-1}, L_h^{t-1} \leftarrow \text{UpdateBounds}(t, \delta, \mathcal{D}_h)$ 
6:   if  $U_1^{t-1}(s_1, c^t) - L_1^{t-1}(s_1, b^t) \leq \varepsilon$  then
7:     return  $b_{t-1}$ , break
8:   end if
9:   // Best
10:   $b^{t-1} \leftarrow \underset{b}{\operatorname{argmin}} [\max_{a \neq b} U_1^{t-1}(s_1, a) - L_1^{t-1}(s_1, b)]$ 
11:  //Challenger
12:   $c^{t-1} \leftarrow \underset{c \neq b^t}{\operatorname{argmax}} U_1^{t-1}(s_1, c)$ 
13:  //Exploration
14:   $a_1^t \leftarrow \underset{a \in \{b^{t-1}, c^{t-1}\}}{\operatorname{argmax}} [U_1^{t-1}(s_1, a) - L_1^{t-1}(s_1, a)]$ 
15:  observe reward  $r_1^t$ , next state  $s_2^t$ , save  $\mathcal{D}_1.append(s_1^t, a_1^t, s_2^t, r_1^t)$ 
16:  for step  $h = 2, \dots, H$  do
17:     $a_h^t \leftarrow \underset{a}{\operatorname{argmax}} U_h^{t-1}(s_h^t, a)$ 
18:    observe reward  $r_{h-1}^t$ , next state  $s_h^t$ , save  $\mathcal{D}_h.append(s_h^t, a_h^t, s_{h+1}^t, r_h^t)$ 
19:  end for
20: end for
```

---





# Result 结果

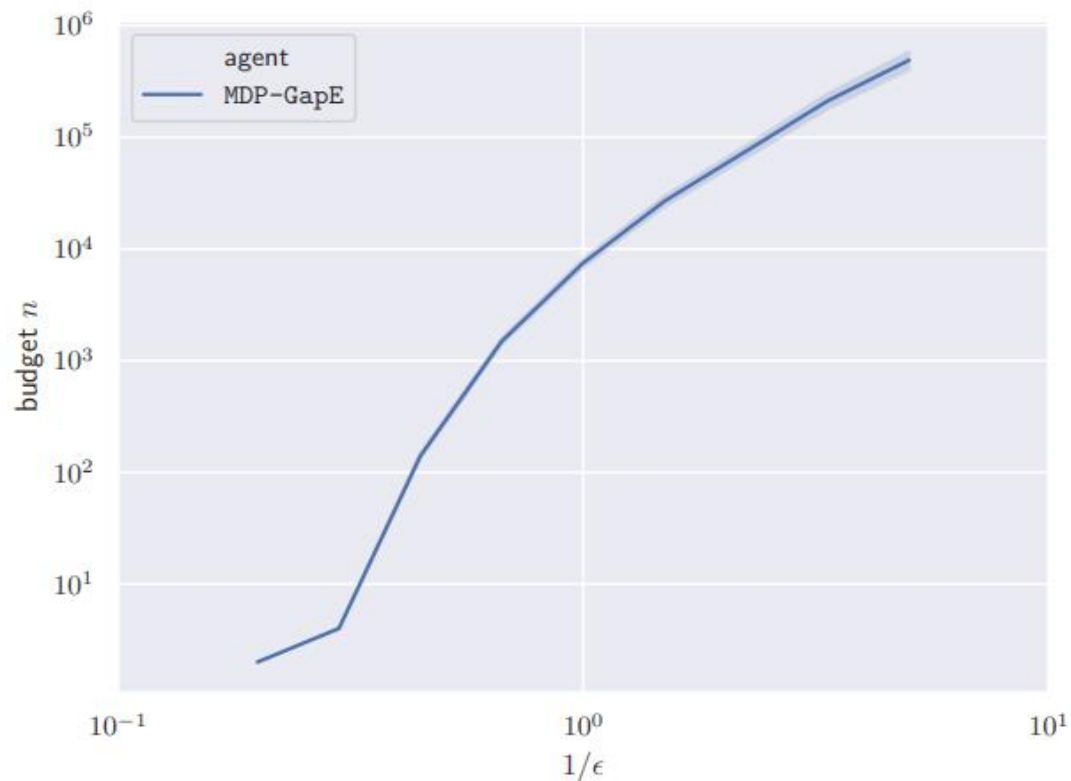


Figure 1: Polynomial dependency of the number  $n$  of oracle calls with respect to  $1/\epsilon$ .

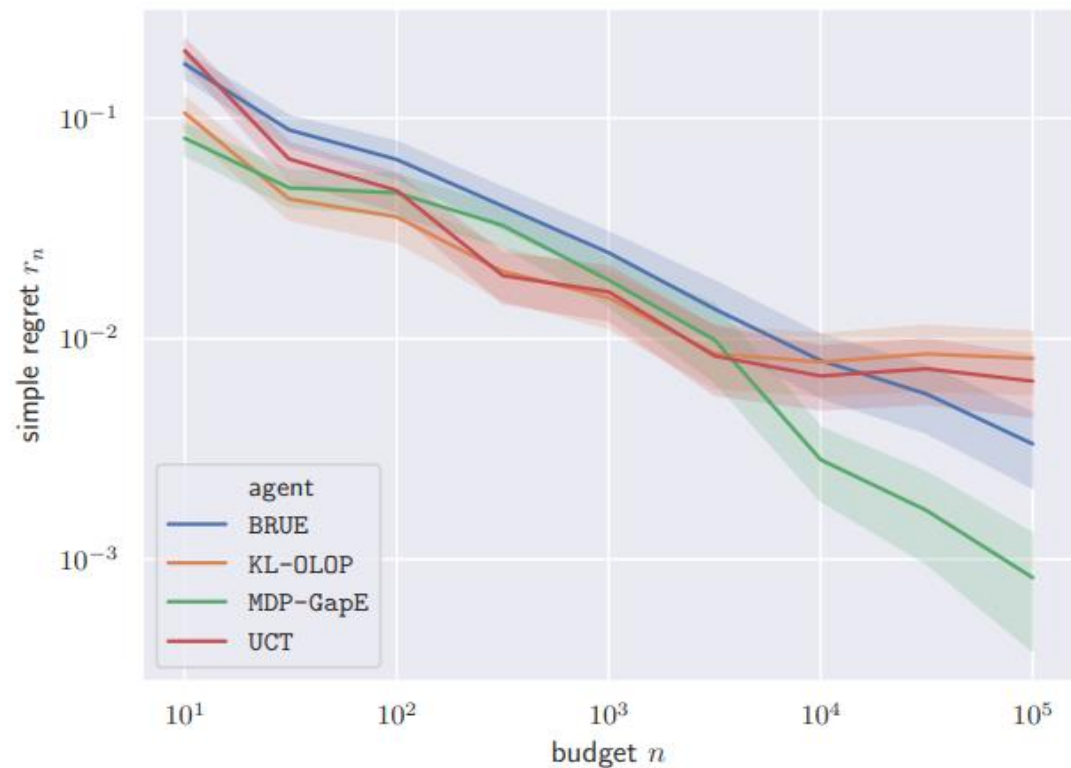


Figure 2: Comparison to KL-OLOP in a fixed-budget setting.